

# Loading and Cleaning Data in R

I know the file exists, why doesn't R?

	River	Site	Ele	Amo	Wea
1	Grasse	Up stream	Al	0.605555555555556	sunny
2	Grasse	Mid stream	Al	0.425	snowy
3	Grasse	Down stream	Al	0.194444444444444	wet
4	Oswegatchie	Up stream	Al	1	cloudy
5	Oswegatchie	Mid stream	Al	0.161111111111111	cloudy
6	Oswegatchie	Down stream	Al	0.0333333333333333	sunny
7	Raquette	Up stream	Al	0.291666666666667	sunny
8	Raquette	Mid stream	Al	0.0388888888888889	cloudy
9	Raquette	Down stream	Al	0	sunny
10	St. Regis	Up stream	Al	0.680555555555556	sunny
11	St. Regis	Mid stream	Al	0.45	snowy
12	St. Regis	Down stream	Al	0.286111111111111	cloudy
13	Grasse	Up stream	Ba	0.505283381364073	wet
14	Grasse	Mid stream	Ba	0.564841498559078	snowy
15	Grasse	Down stream	Ba	0.523535062439962	cloudy
16	Oswegatchie	Up stream	Ba	0.357348703170029	snowy
17	Oswegatchie	Mid stream	Ba	0.560038424591739	sunny
18	Oswegatchie	Down stream	Ba	1	wet
19	Raquette	Up stream	Ba	0	cloudy
20	Raquette	Mid stream	Ba	0.22478386167147	sunny
21	Raquette	Dow stream	Ba	0.364073006724304	cloudy
22	St. Regis	Up stream	Ba	0.379442843419789	wet
23	St. Regis	Mid stream	Ba	0.296829971181556	snowy
24	St. Regis	Down stream	Ba	0.577329490874159	snowy
25	Grasse	Up stream	Br	0.107142857142857	snowy

# R base vs. tidyverse

## R base

- R base is basic R
- Most packages used are installed and loaded by default

# R base vs. **tidyverse**

## R base

- R base is basic R
- Most packages used are installed and loaded by default

## **tidyverse**

- Collection of 'new' packages developed by a team closely affiliated with RStudio
- Packages designed to work well together
- Use a slightly different syntax
- Among others, includes packages used for data transformations and visualizations:
  - e.g., **ggplot2**, **dplyr**, **tidyr**, **readr**

# R base vs. **tidyverse**

## R base

- R base is basic R
- Most packages used are installed and loaded by default

## **tidyverse**

- Collection of 'new' packages developed by a team closely affiliated with RStudio
- Packages designed to work well together
- Use a slightly different syntax
- Among others, includes packages used for data transformations and visualizations:
  - e.g., **ggplot2**, **dplyr**, **tidyr**, **readr**

Can be helpful to understand whether functions are **tidyverse** or R base functions

# Dealing with data

## 1. Loading data

- Get your data into R

## 2. Looking for problems

- Typos
- Incorrectly loaded data

## 3. Fixing problems

- Corrections
- Renaming

## 4. Setting formats

- Dates
- Numbers
- Factors

## 5. Saving your data

# 1. Loading Data

# Data types: What kind of data do you have?

## Specific program files

Type	Extension	R Package	R function
Excel	.xls, .xlsx	<b>readxl</b>	<b>read_excel(sheet = 1)</b>
Open Document	.ods	<b>readODS</b>	<b>read_ods()</b>
SPSS	.sav, .zsav, .por	<b>haven</b>	<b>read_spss()</b>
SAS	.sas7bdat	<b>haven</b>	<b>read_sas()</b>
Stata	.dta	<b>haven</b>	<b>read_dta()</b>
Database Files	.dbf	<b>foreign</b>	<b>read.dbf()</b>

# Data types: What kind of data do you have?

## Specific program files

Type	Extension	R Package	R function
Excel	.xls, .xlsx	<b>readxl</b>	<b>read_excel(sheet = 1)</b>
Open Document	.ods	<b>readODS</b>	<b>read_ods()</b>
SPSS	.sav, .zsav, .por	<b>haven</b>	<b>read_spss()</b>
SAS	.sas7bdat	<b>haven</b>	<b>read_sas()</b>
Stata	.dta	<b>haven</b>	<b>read_dta()</b>
Database Files	.dbf	<b>foreign</b>	<b>read.dbf()</b>

## Convenient but...

- Can be unreliable
- Can take longer

For files that don't change, better to save as a **\*.csv**  
(Comma-separated-variables file)



# Data types: What kind of data do you have?

## General text files

Type	R base	readr package (tidyverse)
Comma separated	<code>read.csv()</code>	<code>read_csv()</code> , <code>read_csv2()</code>
Tab separated	<code>read.delim()</code>	<code>read_tsv()</code>
Space separated	<code>read.table()</code>	<code>read_table()</code>
Fixed-width	<code>read.fwf()</code>	<code>read_fwf()</code>

# Data types: What kind of data do you have?

## General text files

Type	R base	<b>readr</b> package ( <b>tidyverse</b> )
Comma separated	<code>read.csv()</code>	<code>read_csv()</code> , <code>read_csv2()</code>
Tab separated	<code>read.delim()</code>	<code>read_tsv()</code>
Space separated	<code>read.table()</code>	<code>read_table()</code>
Fixed-width	<code>read.fwf()</code>	<code>read_fwf()</code>

- **readr** package especially useful for big data sets (fast!), but have different arguments
- Error/warnings from **readr** are a bit more helpful

# Data types: What kind of data do you have?

## General text files

Type	R base	<b>readr</b> package ( <b>tidyverse</b> )
Comma separated	<code>read.csv()</code>	<code>read_csv()</code> , <code>read_csv2()</code>
Tab separated	<code>read.delim()</code>	<code>read_tsv()</code>
Space separated	<code>read.table()</code>	<code>read_table()</code>
Fixed-width	<code>read.fwf()</code>	<code>read_fwf()</code>

- **readr** package especially useful for big data sets (fast!), but have different arguments
- Error/warnings from **readr** are a bit more helpful

We'll focus on:

- **readxl** package - `read_excel()`
- **readr** package - `read_csv()`, `read_tsv()`

# Where is my data?

```
my_data <- read_csv("weather.csv")
```

```
## Error: 'weather.csv' does not exist in current working directory ('/home/steffi/Projects/Teaching/R  
Workshop/Lessons').
```

With no folder (just file name) R expects file to be in **Working directory**

# Where is my data?

```
my_data <- read_csv("weather.csv")
```

```
## Error: 'weather.csv' does not exist in current working directory ('/home/steffi/Projects/Teaching/R  
Workshop/Lessons').
```

With no folder (just file name) R expects file to be in **Working directory**

## Working directory is:

- Where your RStudio project is
- Your home directory (My Documents, etc.) [If not using RStudio Projects]
- Where you've set it (using `setwd()` or RStudio's Session > Set Working Directory)

# Where is my data?

```
my_data <- read_csv("weather.csv")
```

```
## Error: 'weather.csv' does not exist in current working directory ('/home/steffi/Projects/Teaching/R  
Workshop/Lessons').
```

With no folder (just file name) R expects file to be in **Working directory**

## Working directory is:

- Where your RStudio project is
- Your home directory (My Documents, etc.) [If not using RStudio Projects]
- Where you've set it (using **setwd()** or RStudio's Session > Set Working Directory)

Using Projects in RStudio is a great idea, try to avoid **setwd()**

# Where is my data?

## A note on file paths (file locations)

```
/home
```

- folders separated by /
- **home** is a folder

# Where is my data?

## A note on file paths (file locations)

```
/home/steffi/
```

- folders separated by /
- **home** and **steffi** are folders
- **steffi** is a folder inside of **home**



# Where is my data?

## A note on file paths (file locations)

```
/home/steffi/Documents/R Projects/mydata.csv
```

- folders separated by /
- **home**, **steffi**, **Documents**, **R Projects** are folders
- **steffi** is inside of **home**, **Documents** is inside of **steffi**, etc.
- **mydata.csv** is a data file inside **R Projects** folder

# Where is my data?

## A note on file paths (file locations)

### Absolute Paths

OS	Path
<b>LINUX</b>	/home/steffi/Documents/R Projects/mydata.csv
<b>WINDOWS</b>	C:/Users/steffi/My Documents/R Projects/mydata.csv
<b>MAC</b>	/users/steffi/Documents/R Projects/mydata.csv

Full location, folders and filename

# Where is my data?

## A note on file paths (file locations)

### Absolute Paths

OS	Path
LINUX	/home/steffi/Documents/R Projects/mydata.csv
WINDOWS	C:/Users/steffi/My Documents/R Projects/mydata.csv
MAC	/users/steffi/Documents/R Projects/mydata.csv

Full location, folders and filename

### Relative Paths

Path	Where to look
./mydata.csv	Here (current directory) (./)
../mydata.csv	Go up one directory (../)
./data/mydata.csv	Stay here (./), go into "data" folder (data/)
../data/mydata.csv	Go up one directory (../), then into "data" folder (data/)

Only include some folders, and filename.  
Use relative symbols (**.**/**.** and **../**)

# Where is my data?

## A note on file paths (file locations)

### Absolute Paths

OS	Path
LINUX	/home/steffi/Documents/R Projects/mydata.csv
WINDOWS	C:/Users/steffi/My Documents/R Projects/mydata.csv
MAC	/users/steffi/Documents/R Projects/mydata.csv

### Relative Paths

Path	Where to look
./mydata.csv	Here (current directory) (./)
../mydata.csv	Go up one directory (../)
./data/mydata.csv	Stay here (./), go into "data" folder (data/)
../data/mydata.csv	Go up one directory (../), then into "data" folder (data/)

With RStudio 'Projects' only need to use **relative** paths

Full location, folders and filename

Only include some folders, and filename.  
Use relative symbols (**.**/**.** and **../**)

# Keep yourself organized

## For simple projects

- Create an 'RStudio Project' for each Project (Chapter, Thesis, etc.)
- Create a specific "Data" folder within each project (one per project)

```
- Prospect Lake Quality          # Project Folder
  - prospect_analysis.R
  - Data                         # Data Folder
    - prospect_data_2017-01-01.csv
    - prospect_data_2017-02-01.csv
```

# Keep yourself organized

## For simple projects

- Create an 'RStudio Project' for each Project (Chapter, Thesis, etc.)
- Create a specific "Data" folder within each project (one per project)

```
- Prospect Lake Quality          # Project Folder
  - prospect_analysis.R
  - Data                        # Data Folder
    - prospect_data_2017-01-01.csv
    - prospect_data_2017-02-01.csv
```

- Use **relative** paths to refer to this folder ("./")

```
d <- read_csv("./data/prospect_data_2017-01-01.csv")
```

**Let's Load Some Data!**

# Your turn: Load some data

1. Save/move the **rivers\_correct.xlsx** file to a "Data" folder in your project (download [here](#))

2. Load the package

```
library(readxl)
```

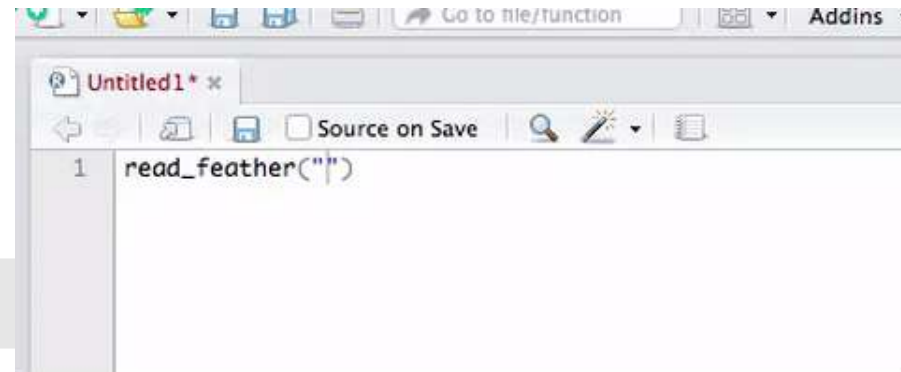
3. Read in the Excel file and assign to object **rivers**

```
rivers <- read_excel("./data/rivers_correct.xlsx")
```

4. Use **head()** and **tail()** functions to look at the data

```
head(rivers)
tail(rivers)
```

5. Click on the data object in your "Environment" pane to look at the whole data set



Use the 'tab' key in RStudio when typing in the file name for auto-complete



# How do I know which function to use?

## Look at the file extension:

- **rivers\_correct.csv** (download [here](#))
- **.csv** = Comma-separated-variables = **read\_csv()**

# How do I know which function to use?

## Look at the file extension:

- **rivers\_correct.csv** (download [here](#))
- **.csv** = Comma-separated-variables = **read\_csv()**

But not always obvious...

# How do I know which function to use?

## Look at the file:

- **master\_moch.txt** (download [here](#))
- In lower right-hand pane, click on **Files**
  - Click on **Data** folder
  - Click on **master\_moch.txt**
  - Click "View File"

ID	region	hab	freq	freq.sd	p.notes
MCB02	kam	0.5266879074	3.9806600009	3.9806600009	0.4592592593
MCB03	kam	-0.9707703735	4.1090031783	4.1090031783	0.5
MCB04	kam	-0.9707703735	4.2463067674	4.2463067674	0.5151515152

This **does not** read the file into R, but only shows you the contents as text.

# How do I know which function to use?

## Look at the file:

- **master\_moch.txt** (download [here](#))
- In lower right-hand pane, click on **Files**
  - Click on **Data** folder
  - Click on **master\_moch.txt**
  - Click "View File"

Hmm, not comma-separated,  
maybe tab-separated?

ID	region	hab	freq	freq.sd	p.notes
MCB02	kam	0.5266879074	3.9806600009	3.9806600009	0.4592592593
MCB03	kam	-0.9707703735	4.1090031783	4.1090031783	0.5
MCB04	kam	-0.9707703735	4.2463067674	4.2463067674	0.5151515152

This **does not** read the file into R, but only shows you the contents as text.

# How do I know what to use?

## Peak:

- Pick a read function with your best guess (`read_csv()` is a good start)
- Use `n_max` to read only first few rows

```
read_csv("./data/master_moch.txt", n_max = 3)
```

```
## # A tibble: 3 x 1
##   `ID\tregion\thab\tfreq\tfreq.sd\tp.notes`
##   <chr>
## 1 "MCB02\tkam\t0.5266879074\t3.9806600009\t3.9806600009\t0.4592592593"
## 2 "MCB03\tkam\t-0.9707703735\t4.1090031783\t4.1090031783\t0.5"
## 3 "MCB04\tkam\t-0.9707703735\t4.2463067674\t4.2463067674\t0.5151515152"
```

`\t` means tab, so this is tab-separated data

# How do I know what to use?

## Peak:

- Try again with `read_tsv()`

```
read_tsv("./data/master_moch.txt", n_max = 3) # note change in function!
```

```
## # A tibble: 3 x 6
```

```
##   ID      region    hab  freq freq.sd p.notes
```

```
##   <chr> <chr>    <dbl> <dbl>   <dbl>   <dbl>
```

```
## 1 MCB02 kam      0.527  3.98    3.98    0.459
```

```
## 2 MCB03 kam     -0.971  4.11    4.11    0.5
```

```
## 3 MCB04 kam     -0.971  4.25    4.25    0.515
```

Excellent!

# Specifics of loading functions

# 1. col\_names

- Geolocator data (download [here](#))

```
my_data <- read_csv("./data/geolocators.csv")  
my_data
```

```
## # A tibble: 20 x 2
```

```
##   `02/05/11 22:29:59` `64`
```

```
##   <chr>                <dbl>
```

```
## 1 02/05/11 22:31:59      64
```

```
## 2 02/05/11 22:33:59      38
```

```
## 3 02/05/11 22:35:59      38
```

```
## 4 02/05/11 22:37:59      34
```

```
## 5 02/05/11 22:39:59      30
```

```
## 6 02/05/11 22:41:59      34
```

```
## 7 02/05/11 22:43:59      40
```

```
## 8 02/05/11 22:45:59      46
```

```
## 9 02/05/11 22:47:59      48
```

```
## 10 02/05/11 22:49:59     46
```

```
## # ... with 10 more rows
```

Oops?



# 1. col\_names

- Geolocator data (download [here](#))

```
my_data <- read_csv("./data/geolocators.csv")  
my_data
```

```
## # A tibble: 20 x 2  
##   `02/05/11 22:29:59` `64`  
##   <chr>                <dbl>  
## 1 02/05/11 22:31:59      64  
## 2 02/05/11 22:33:59      38  
## 3 02/05/11 22:35:59      38  
## 4 02/05/11 22:37:59      34  
## 5 02/05/11 22:39:59      30  
## 6 02/05/11 22:41:59      34  
## 7 02/05/11 22:43:59      40  
## 8 02/05/11 22:45:59      46  
## 9 02/05/11 22:47:59      48  
## 10 02/05/11 22:49:59      46  
## # ... with 10 more rows
```

- **read\_csv**, **read\_tsv**, etc. assume that the first row contains the column names
- This file doesn't have headers

Oops?

# 1. col\_names

## Declare no headings

```
my_data <- read_csv("./data/geolocators.csv",  
                    col_names = FALSE)  
my_data
```

```
## # A tibble: 21 x 2
```

```
##      X1                X2
```

```
##      <chr>            <dbl>
```

```
##  1 02/05/11 22:29:59    64
```

```
##  2 02/05/11 22:31:59    64
```

```
##  3 02/05/11 22:33:59    38
```

```
##  4 02/05/11 22:35:59    38
```

```
##  5 02/05/11 22:37:59    34
```

```
##  6 02/05/11 22:39:59    30
```

```
##  7 02/05/11 22:41:59    34
```

```
##  8 02/05/11 22:43:59    40
```

```
##  9 02/05/11 22:45:59    46
```

```
## 10 02/05/11 22:47:59    48
```

```
## # ... with 11 more rows
```

# 1. col\_names

## Declare no headings

```
my_data <- read_csv("./data/geolocators.csv",  
                    col_names = FALSE)  
my_data
```

```
## # A tibble: 21 x 2
```

##	X1	X2
##	<chr>	<dbl>
##	1 02/05/11 22:29:59	64
##	2 02/05/11 22:31:59	64
##	3 02/05/11 22:33:59	38
##	4 02/05/11 22:35:59	38
##	5 02/05/11 22:37:59	34
##	6 02/05/11 22:39:59	30
##	7 02/05/11 22:41:59	34
##	8 02/05/11 22:43:59	40
##	9 02/05/11 22:45:59	46
##	10 02/05/11 22:47:59	48
##	# ... with 11 more rows	

## Name headings

```
my_data <- read_csv("./data/geolocators.csv",  
                    col_names = c("date", "light"))  
my_data
```

```
## # A tibble: 21 x 2
```

##	date	light
##	<chr>	<dbl>
##	1 02/05/11 22:29:59	64
##	2 02/05/11 22:31:59	64
##	3 02/05/11 22:33:59	38
##	4 02/05/11 22:35:59	38
##	5 02/05/11 22:37:59	34
##	6 02/05/11 22:39:59	30
##	7 02/05/11 22:41:59	34
##	8 02/05/11 22:43:59	40
##	9 02/05/11 22:45:59	46
##	10 02/05/11 22:47:59	48
##	# ... with 11 more rows	

## 2. skip info rows before data

- Grain size data (download [here](#))

```
my_data <- read_tsv("./data/grain_size.txt")
```

```
## Warning: Missing column names filled in: 'X2' [2], 'X3' [3], 'X4' [4], 'X5' [5], 'X6' [6], 'X7' [7]
```

```
my_data
```

```
## # A tibble: 36 x 7
```

##	`DATA DOWNLOAD: 2015-09-23`	X2	X3	X4	X5	X6	X7
##	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
##	1 SYSTEM 001	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>
##	2 LOGGER X	<NA>	<NA>	<NA>	<NA>	<NA>	<NA>
##	3 lab_num	CSP	sample_num	depth_lb	csa	msa	fsa
##	4 3177	CSP01	CSP01-P-1-1	4	13.04	17.37	8.19
##	5 3178	CSP01	CSP01-P-1-2	12	10.74	16.9	7.92
##	6 3179	CSP01	CSP01-P-1-3	35	12.11	17.75	6.99
##	7 3180	CSP01	CSP01-P-1-4	53	17.61	18.16	6.29
##	8 3181	CSP01	CSP01-P-1-5	83	21.05	18.38	6.26
##	9 3182	CSP01	CSP01-P-1-6	105	19.02	18.43	6.28

## 2. **skip** info rows before data

- Grain size data (download [here](#))

```
my_data <- read_tsv("./data/grain_size.txt")
```

### Look at the file:

- Go to **Files** tab (lower right-hand pane)
- Click on **Data** folder
- Click on **grain\_size.txt** file
- Click "**View file**"

## 2. skip info rows before data

- Grain size data (download [here](#))

```
my_data <- read_tsv("./data/grain_size.txt")
```

### Look at the file:

- Go to **Files** tab (lower right-hand pane)
- Click on **Data** folder
- Click on **grain\_size.txt** file
- Click "**View file**"

DATA DOWNLOAD: 2015-09-23

SYSTEM 001

LOGGER X

lab_num	CSP	sample_num	depth_lb	csa	msa	fsa
3177	CSP01	CSP01-P-1-1	4	13.04	17.37	8.19
3178	CSP01	CSP01-P-1-2	12	10.74	16.9	7.92
3179	CSP01	CSP01-P-1-3	35	12.11	17.75	6.99
3180	CSP01	CSP01-P-1-4	53	17.61	18.16	6.29
3181	CSP01	CSP01-P-1-5	83	21.05	18.38	6.26

**Ah ha!**

Metadata was stored at the top of the file

## 2. **skip** info rows before data

- Grain size data (download [here](#))
- Add **skip = 3** to skip the first three rows

```
my_data <- read_tsv("./data/grain_size.txt", skip = 3)
my_data
```

```
## # A tibble: 33 x 7
##   lab_num CSP   sample_num depth_lb   csa   msa   fsa
##   <dbl> <chr> <chr>         <dbl> <dbl> <dbl> <dbl>
## 1   3177 CSP01 CSP01-P-1-1         4 13.0  17.4  8.19
## 2   3178 CSP01 CSP01-P-1-2        12 10.7  16.9  7.92
## 3   3179 CSP01 CSP01-P-1-3        35 12.1  17.8  6.99
## 4   3180 CSP01 CSP01-P-1-4        53 17.6  18.2  6.29
## 5   3181 CSP01 CSP01-P-1-5        83 21.0  18.4  6.26
## 6   3182 CSP01 CSP01-P-1-6       105 19.0  18.4  6.28
## 7   3183 CSP08 CSP08-P-1-1        10 11.6  17.1  8.18
## 8   3184 CSP08 CSP08-P-1-2        27 15.4  16.2  6.76
## 9   3185 CSP08 CSP08-P-1-3        90 14.9  15.8  7.12
## 10  3186 CSP02 CSP02-P-1-1         5  8.75  8.64  3.41
## # ... with 23 more rows
```

# Your turn: Load this data set

Try loading the telemetry data set: **Sta A Data 2006-11-07.dmp** (download it [here](#))

1. Look at the file (click on the file in your File window in RStudio)
2. Decide function based on the delimiter (comma, space, or tab?)
3. Any other options need to be specified?

**Extra Challenge**  
Load some of your own  
tricky data

## It should look like this:

```
## # A tibble: 19 x 7
##   StartDate Time      Frequency `Rate/Temp`    Pwr Ant      SD
##   <dbl> <time>      <dbl>      <dbl> <dbl> <chr> <dbl>
## 1 39022 17:15:36    150.      34.8  175 M0      0
## 2 39022 17:19:14    148.      19.2   72 M0      0
## 3 39022 17:19:25    148.      19.7  194 M1      0
## 4 39022 17:20:04    149.      33.8  104 M0      0
## 5 39022 17:20:17    149.      33.7  152 M1      0
## 6 39022 17:20:57    150.      34.2  188 M0      0
## 7 39022 17:22:50    148.       9.8  188 M0      0
## # ... with 12 more rows
```



## 2. Looking for problems

# Look at the data

- Make sure columns as expected (correctly assigned file format)
- Make sure no extra lines above the data (should we have used a skip?)
- Make sure column names look appropriate

```
library(palmerpenguins)
penguins
```

```
## # A tibble: 344 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex   year
##   <fct>   <fct>         <dbl>         <dbl>             <int>         <int> <fct> <int>
## 1 Adelie  Torgersen         39.1          18.7             181          3750 male   2007
## 2 Adelie  Torgersen         39.5          17.4             186          3800 female 2007
## 3 Adelie  Torgersen         40.3           18             195          3250 female 2007
## 4 Adelie  Torgersen          NA           NA              NA           NA <NA>   2007
## 5 Adelie  Torgersen         36.7          19.3             193          3450 female 2007
## 6 Adelie  Torgersen         39.3          20.6             190          3650 male   2007
## 7 Adelie  Torgersen         38.9          17.8             181          3625 female 2007
## 8 Adelie  Torgersen         39.2          19.6             195          4675 male   2007
## 9 Adelie  Torgersen         34.1          18.1             193          3475 <NA>   2007
## 10 Adelie Torgersen         42           20.2             190          4250 <NA>   2007
## # ... with 334 more rows
```

# Look at the data

- Did the whole data set load?
- Are there extra blank lines at the end of the data?

```
tail(penguins)
```

```
## # A tibble: 6 x 8
```

##	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex	year
##	<fct>	<fct>	<dbl>	<dbl>	<int>	<int>	<fct>	<int>
## 1	Chinstrap	Dream	45.7	17	195	3650	female	2009
## 2	Chinstrap	Dream	55.8	19.8	207	4000	male	2009
## 3	Chinstrap	Dream	43.5	18.1	202	3400	female	2009
## 4	Chinstrap	Dream	49.6	18.2	193	3775	male	2009
## 5	Chinstrap	Dream	50.8	19	210	4100	male	2009
## 6	Chinstrap	Dream	50.2	18.7	198	3775	female	2009

# skim()

- Are the columns appropriate?
- Are numeric values appropriate? Should there be NAs?
- Are there any typos in the factors?
- Are the columns in the appropriate format (i.e., numeric, character, factor, date)
- Are there as many observations as you expected?

```
library(skimr)
skim(penguins)
```

```
## — Data Summary —————
##                               Values
## Name                         penguins
## Number of rows                344
## Number of columns             8
## -----
## Column type frequency:
##   factor                       3
##   numeric                      5
## -----
## Group variables               None
```

# skim()

```
##
## — Variable type: factor

##   skim_variable n_missing complete_rate ordered n_unique top_counts
## 1 species           0           1      FALSE           3 Ade: 152, Gen: 124, Chi: 68
## 2 island            0           1      FALSE           3 Bis: 168, Dre: 124, Tor: 52
## 3 sex               11          0.968 FALSE           2 mal: 168, fem: 165
```

```
##
```

```
## — Variable type: numeric
```

```
##   skim_variable  n_missing complete_rate  mean    sd    p0    p25    p50    p75    p100 hist
## 1 bill_length_mm      2         0.994  43.9   5.46   32.1   39.2   44.4   48.5   59.6  ████
## 2 bill_depth_mm      2         0.994  17.2   1.97   13.1   15.6   17.3   18.7   21.5  ████
## 3 flipper_length_mm  2         0.994  201.   14.1   172    190    197    213    231   ████
## 4 body_mass_g        2         0.994 4202.  802.   2700   3550   4050   4750   6300  ████
## 5 year               0           1    2008.  0.818 2007   2007   2008   2009   2009  ████
```

# count()

- Check for typos in categorical columns

```
count(penguins, species)
```

```
## # A tibble: 3 x 2
##   species      n
##   <fct>    <int>
## 1 Adelie   152
## 2 Chinstrap 68
## 3 Gentoo   124
```

```
count(penguins, island)
```

```
## # A tibble: 3 x 2
##   island      n
##   <fct>    <int>
## 1 Biscoe   168
## 2 Dream    124
## 3 Torgersen 52
```

# Example of problematic data

```
rivers <- read_csv("../data/rivers_correct.csv")
rivers
```

```
## # A tibble: 300 x 5
##   `River Name` Site      Ele    Amo           Wea
##   <chr>         <chr>    <chr> <chr>         <chr>
## 1 Grasse       Up stream Al    0.605555555555556 sunny
## 2 Grasse       Mid stream Al    0.425          wet
## 3 Grase        Down stream Al    0.194444444444444 wet
## 4 Oswegatchie Up stream  Al    1             cloudy
## 5 Oswegatchie Mid stream Al    0.161111111111111 sunny
## 6 Oswegatchie Down stream Al    0.033333333333333 sunny
## 7 Raquette     Up stream  Al    0.291666666666667 sunny
## 8 Raquette     Mid stream Al    0.038888888888889 sunny
## 9 Raquette     Down stream Al    0             snowy
## 10 St. Regis   Up stream  Al    0.680555555555556 cloudy
## # ... with 290 more rows
```

- Column names are not all clean (**River Name** or obvious (what is **Ele**?)
- **Amo** - should be numeric but isn't
- At least one typo in River (**Grase** should be **Grasse**)

# Example of problematic data

```
skim(rivers)
```

```
##
```

```
## — Variable type: character
```

##	skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
## 1	River Name	0	1	5	11	0	8	0
## 2	Site	0	1	8	11	0	5	0
## 3	Ele	0	1	1	2	0	25	0
## 4	Amo	12	0.96	1	19	0	198	0
## 5	Wea	0	1	3	6	0	4	0

- Not much additional info here



# Example of problematic data

```
count(rivers, `River Name`)
```

```
## # A tibble: 8 x 2
##   `River Name`      n
##   <chr>          <int>
## 1 Grase           1
## 2 Grass           1
## 3 grasse          1
## 4 Grasse         72
## 5 Oswegatchie     75
## 6 raquette        1
## 7 Raquette        74
## 8 St. Regis       75
```

```
count(rivers, Site)
```

```
## # A tibble: 5 x 2
##   Site              n
##   <chr>          <int>
## 1 Dow stream       1
## 2 Down stream     99
## 3 Mid stream     100
## 4 Up stream       99
## 5 Upstream        1
```

Typos in both categorical columns

### **3. Fixing problems**

# Cleaning column names

## `clean_names()`

```
library(janitor)
rivers <- clean_names(rivers)
rivers
```

```
## # A tibble: 300 x 5
##   river_name site      ele amo      wea
##   <chr>      <chr>    <chr> <chr>    <chr>
## 1 Grasse    Up stream  Al    0.605555555555556 sunny
## 2 Grasse    Mid stream Al    0.425      wet
## 3 Grasse    Down stream Al    0.194444444444444 wet
## 4 Oswegatchie Up stream  Al    1          cloudy
## 5 Oswegatchie Mid stream Al    0.161111111111111 sunny
## 6 Oswegatchie Down stream Al    0.0333333333333333 sunny
## 7 Raquette   Up stream  Al    0.291666666666667 sunny
## 8 Raquette   Mid stream Al    0.0388888888888889 sunny
## 9 Raquette   Down stream Al    0          snowy
## 10 St. Regis  Up stream  Al    0.680555555555556 cloudy
## # ... with 290 more rows
```

# Cleaning column names

## `rename()` columns

```
rivers <- rename(rivers, element = ele, amount = amo)
rivers
```

```
## # A tibble: 300 x 5
##   river_name site      element amount      wea
##   <chr>      <chr>      <chr>  <chr>      <chr>
## 1 Grasse     Up stream  Al      0.605555555555556 sunny
## 2 Grasse     Mid stream Al      0.425      wet
## 3 Grasse     Down stream Al      0.194444444444444 wet
## 4 Oswegatchie Up stream  Al      1          cloudy
## 5 Oswegatchie Mid stream Al      0.161111111111111 sunny
## 6 Oswegatchie Down stream Al      0.0333333333333333 sunny
## 7 Raquette   Up stream  Al      0.291666666666667 sunny
## 8 Raquette   Mid stream Al      0.0388888888888889 sunny
## 9 Raquette   Down stream Al      0          snowy
## 10 St. Regis  Up stream  Al      0.680555555555556 cloudy
## # ... with 290 more rows
```

# Subsetting columns

## **select()** columns you do want

```
rivers <- select(rivers, river_name, site, element, amount)
```

# Subsetting columns

## `select()` columns you do want

```
rivers <- select(rivers, river_name, site, element, amount)
```

## OR, `unselect()` columns you don't want

```
rivers <- select(rivers, -wea)  
rivers
```

```
## # A tibble: 300 x 4  
##   river_name site      element amount  
##   <chr>      <chr>      <chr>    <chr>  
## 1 Grasse    Up stream  Al      0.6055555555555556  
## 2 Grasse    Mid stream Al      0.425  
## 3 Grasse    Down stream Al      0.1944444444444444  
## 4 Oswegatchie Up stream  Al      1  
## 5 Oswegatchie Mid stream Al      0.1611111111111111  
## 6 Oswegatchie Down stream Al      0.03333333333333333  
## 7 Raquette   Up stream  Al      0.2916666666666667  
## 8 Raquette   Mid stream Al      0.03888888888888889
```

# Cleaning columns

## Put it all together

```
rivers <- read_csv("../data/rivers_correct.csv")
rivers <- clean_names(rivers)
rivers <- rename(rivers, element = ele, amount = amo)
rivers <- select(rivers, -wea)
rivers
```

```
## # A tibble: 300 x 4
##   river_name site      element amount
##   <chr>      <chr>      <chr>    <chr>
## 1 Grasse    Up stream  Al      0.6055555555555556
## 2 Grasse    Mid stream Al      0.425
## 3 Grasse    Down stream Al      0.1944444444444444
## 4 Oswegatchie Up stream  Al      1
## 5 Oswegatchie Mid stream Al      0.1611111111111111
## 6 Oswegatchie Down stream Al      0.0333333333333333
## 7 Raquette   Up stream  Al      0.2916666666666667
## 8 Raquette   Mid stream Al      0.0388888888888889
## 9 Raquette   Down stream Al      0
```

# Cleaning columns

## Put it all together

```
rivers <- read_csv("./data/rivers_correct.csv")
rivers <- clean_names(rivers)
rivers <- rename(rivers, element = ele, amount = amo)
rivers <- select(rivers, -wea)
rivers
```

Note repeated data frame  
**rivers**


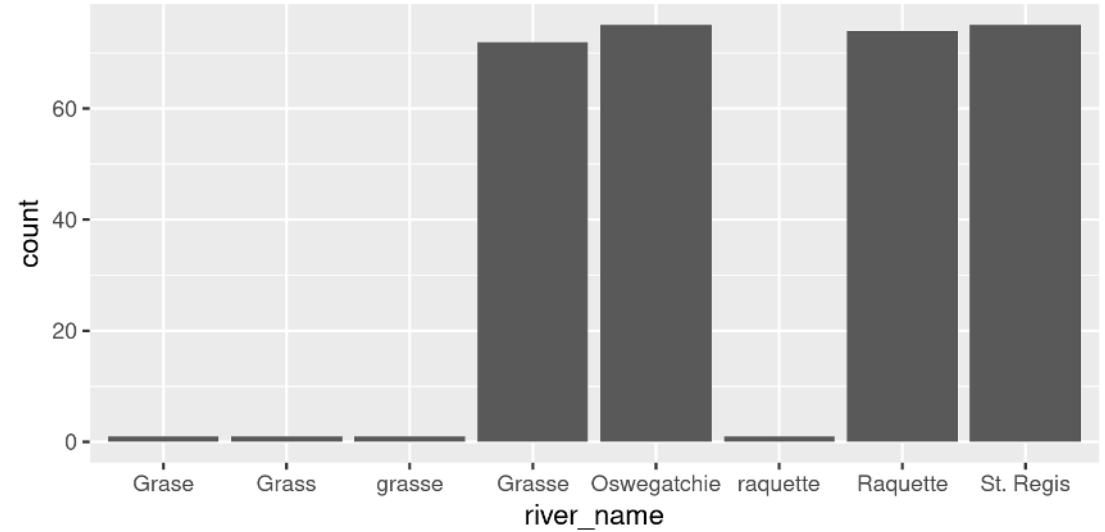
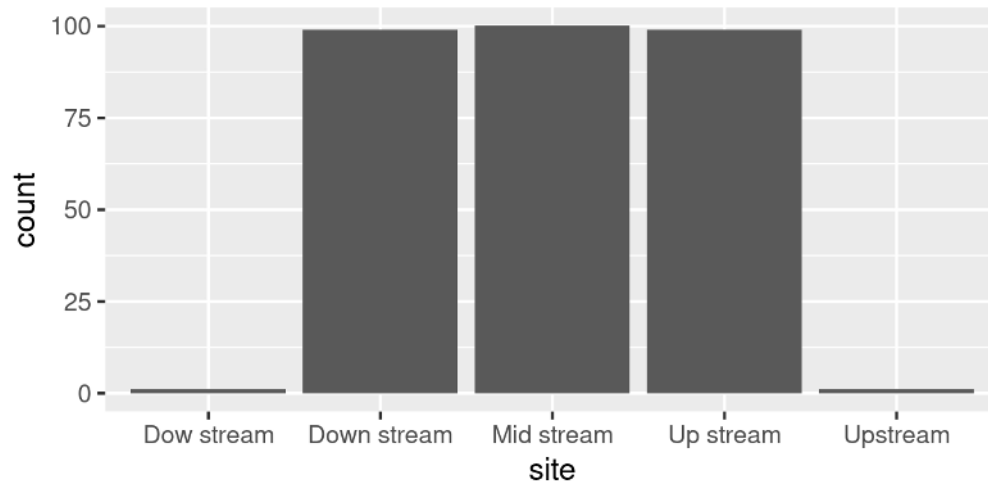
```
## # A tibble: 300 x 4
##   river_name site      element amount
##   <chr>      <chr>      <chr>    <chr>
## 1 Grasse    Up stream  Al      0.6055555555555556
## 2 Grasse    Mid stream Al      0.425
## 3 Grasse    Down stream Al      0.1944444444444444
## 4 Oswegatchie Up stream  Al      1
## 5 Oswegatchie Mid stream Al      0.1611111111111111
## 6 Oswegatchie Down stream Al      0.03333333333333333
## 7 Raquette   Up stream  Al      0.2916666666666667
## 8 Raquette   Mid stream Al      0.03888888888888889
## 9 Raquette   Down stream Al      0
```



# Fixing typos

## Look for typos (Visually)

```
ggplot(data = rivers, aes(x = river_name)) + geom_bar()
ggplot(data = rivers, aes(x = site)) + geom_bar()
ggplot(data = rivers, aes(x = element)) + geom_bar()
```



# Fixing typos

## Look for typos with `count()`

```
count(rivers, river_name)
```

```
## # A tibble: 8 x 2
##   river_name      n
##   <chr>      <int>
## 1 Grase         1
## 2 Grass         1
## 3 grasse        1
## 4 Grasse       72
## 5 Oswegatchie  75
## 6 raquette      1
## 7 Raquette     74
## 8 St. Regis    75
```

## `filter()` the data to highlight them

```
filter(rivers, river_name == "Grase")
```

```
## # A tibble: 1 x 4
##   river_name site      element amount
##   <chr>      <chr>      <chr>   <chr>
## 1 Grase     Down stream Al      0.19444444444444444
```

# Fixing typos

## Replace typos

Combine the **if\_else()** / **case\_when()** functions with **mutate()** function

**mutate()** creates or changes columns in a data frame:

```
mutate(dataframe, column = new_values)
```

**if\_else()** tests for a condition, and returns one value if FALSE and another if TRUE

```
if_else(condition, value_if_true, value_if_false)
```

**case\_when()** tests for multiple conditions, and returns different values depending

```
case_when(condition1 ~ value_if_true,  
          condition2 ~ value_if_true,  
          condition3 ~ value_if_true,  
          TRUE ~ default_value)
```

# Fixing typos

## Replace typos

Combine the **if\_else** function with the **mutate()** function

```
rivers <- mutate(rivers, river_name = if_else(river_name == "Grase", "Grasse", river_name))
```

## Check that it's gone:

```
filter(rivers, river_name == "Grase")
```

```
## # A tibble: 0 x 4
```

```
## # ... with 4 variables: river_name <chr>, site <chr>, element <chr>, amount <chr>
```

# Iterative process

- Make some corrections
- Check the data
- Make some more corrections (either add to or modify existing code)

# Your Turn: Fix another one of the "Grasse" typos

1. Check the data with **count()**
2. Use **mutate()** and **replace()** to fix the typo

**Extra Challenge**  
Examine and fix problems  
in your own data

```
rivers <- read_csv("./data/rivers_correct.csv")
rivers <- clean_names(rivers)
rivers <- rename(rivers, element = ele, amount = amo)
rivers <- select(rivers, -wea)
rivers <- mutate(rivers, river_name = if_else(river_name == "Grase", "Grasse", river_name))
```

```
rivers <- mutate(???, ??? = ???)
```

# Fixing typos

## To be more efficient, fix all typos at once

```
rivers <- read_csv("./data/rivers_correct.csv")
rivers <- clean_names(rivers)
rivers <- rename(rivers, element = ele, amount = amo)
rivers <- select(rivers, -wea)
rivers <- mutate(rivers, river_name = if_else(river_name %in% c("Grase", "Grass", "grasse"),
                                              "Grasse",
                                              river_name))
```

**==** compares one item to one other

**%in%** compares one item to many  
different ones

# Tangent: Why use tidyverse packages?

## Pipes! `%>%` Allow you to string commands together

Instead of:

```
rivers <- read_csv("./data/rivers_correct.csv")
rivers <- clean_names(rivers)
rivers <- rename(rivers, element = ele, amount = amo)
rivers <- select(rivers, -wea)
rivers <- mutate(rivers,
  river_name = if_else(river_name == "Grase", "Grasse", river_name),
  site = if_else(site == "Dow stream", "Down stream", site))
```



# Tangent: Why use tidyverse packages?

## Pipes! `%>%` Allow you to string commands together

Instead of:

```
rivers <- read_csv("../data/rivers_correct.csv")
rivers <- clean_names(rivers)
rivers <- rename(rivers, element = ele, amount = amo)
rivers <- select(rivers, -wea)
rivers <- mutate(rivers,
  river_name = if_else(river_name == "Grase", "Grasse", river_name),
  site = if_else(site == "Dow stream", "Down stream", site))
```

We have:

```
rivers <- read_csv("../data/rivers_correct.csv") %>%
  clean_names() %>%
  rename(element = ele, amount = amo) %>%
  select(-wea) %>%
  mutate(river_name = if_else(river_name == "Grase", "Grasse", river_name),
    site = if_else(site == "Dow stream", "Down stream", site))
```

# Play around

Take a moment to play with this code in your console

## Convert this:

```
rivers <- read_csv("./data/rivers_correct.csv")
rivers <- clean_names(rivers)
rivers <- rename(rivers, element = ele, amount = amo)
rivers <- select(rivers, -wea)
rivers <- mutate(rivers,
  river_name = if_else(river_name %in% c("Grase", "Grass", "grasse"), "Grasse", river_name),
  site = if_else(site == "Dow stream", "Down stream", river_name))
```

## To this:

```
rivers <- read_csv("./data/rivers_correct.csv") %>%
  clean_names() %>%
  rename(element = ele, amount = amo) %>%
  select(-wea) %>%
  mutate(river_name = if_else(river_name %in% c("Grase", "Grass", "grasse"), "Grasse", river_name),
    site = if_else(site == "Dow stream", "Down stream", site))
```

# Your turn: Fix the remaining typos

- Remember this is an iterative process (you may find your self reloading the data often)
- Don't worry about numerical problems for now

## Add to this code

```
rivers <- read_csv("./data/rivers_correct.csv") %>%  
  clean_names() %>%  
  rename(element = ele, amount = amo) %>%  
  select(-wea) %>%  
  mutate(river_name = if_else(river_name %in% c("Grase", "Grass", "grasse"), "Grasse", river_name),  
         site = if_else(site == "Dow stream", "Down stream", site))
```

Expect **river\_name**: Grasse, Oswegatchie, Raquette, St.Regis

Expect **site**: Down stream, Up stream

**Extra Challenge**  
Examine and fix problems  
in your own data

## Comparing single items

```
A == "hello"  
A %in% "hello"
```

## Comparing multiple items

```
A %in% c("hello", "bye")  
# NOT A == c("hello", "bye")
```

# Your turn: Fix the remaining typos

- Remember this is an iterative process (you may find your self reloading the data often)
- Don't worry about numerical problems for now

## Add to this code

```
rivers <- read_csv("./data/rivers_correct.csv") %>%  
  clean_names() %>%  
  rename(element = ele, amount = amo) %>%  
  select(-wea) %>%  
  mutate(river_name = if_else(river_name %in% c("Grase", "Grass", "grasse"), "Grasse", river_name),  
         site = if_else(site == "Dow stream", "Down stream", site))
```

Expect **river\_name**: Grasse, Oswegatchie, Raquette, St.Regis

Expect **site**: Down stream, Up stream

**Extra Challenge**  
Examine and fix problems  
in your own data

## Comparing single items

```
A == "hello"  
A %in% "hello"
```

## Comparing multiple items

```
A %in% c("hello", "bye")  
# NOT A == c("hello", "bye")
```

## 4. Fixing formats

# Typos that affect classes (formats)

## Look for problems

```
rivers
```

```
## # A tibble: 300 x 4
##   river_name site      element amount
##   <chr>      <chr>      <chr>    <chr>
## 1 Grasse     Up stream  Al       0.6055555555555556
## 2 Grasse     Mid stream Al       0.425
## 3 Grasse     Down stream Al      0.1944444444444444
## 4 Oswegatchie Up stream  Al       1
## 5 Oswegatchie Mid stream Al      0.1611111111111111
## 6 Oswegatchie Down stream Al      0.03333333333333333
## 7 Raquette   Up stream  Al      0.2916666666666667
## 8 Raquette   Mid stream Al      0.03888888888888889
## 9 Raquette   Down stream Al       0
## 10 St. Regis Up stream  Al      0.6805555555555556
## # ... with 290 more rows
```

Why all character (chr)?

# Changing classes

Function	Input	Output
<code>as.character()</code>	Any vector	Text (Characters)
<code>as.numeric()</code>	Any vector (but returns NAs if not numbers)	Numbers
<code>as.logical()</code>	TRUE, FALSE, T, F, 0 (FALSE), any other number (all TRUE)	TRUE or FALSE
<code>as.factor()</code>	Any vector	Categories

# Changing classes

Function	Input	Output
<code>as.character()</code>	Any vector	Text (Characters)
<code>as.numeric()</code>	Any vector (but returns NAs if not numbers)	Numbers
<code>as.logical()</code>	TRUE, FALSE, T, F, 0 (FALSE), any other number (all TRUE)	TRUE or FALSE
<code>as.factor()</code>	Any vector	Categories

## For example...

```
a <- c(1, 2, 10)
```

```
as.character(a)
```

```
## [1] "1" "2" "10"
```

```
as.numeric(a)
```

```
## [1] 1 2 10
```

```
b <- c("hello", "bye", 1)
```

```
as.character(b)
```

```
## [1] "hello" "bye" "1"
```

```
as.numeric(b)
```

```
## Warning: NAs introduced by coercion
```

```
## [1] NA NA 1
```

We'll deal with dates and times later...



# Fixing numerical typos

## Find the problem (when we don't know what they are)

- Make a new column and convert **amount** to numbers

```
rivers <- mutate(rivers, amount2 = as.numeric(amount))
```

```
## Warning: Problem with `mutate()` input `amount2`.
```

```
## i NAs introduced by coercion
```

```
## i Input `amount2` is `as.numeric(amount)`.
```

```
## Warning in mask$eval_all_mutate(dots[[i]]): NAs introduced by coercion
```

This warning tells us that some values didn't convert to numbers

# Fixing numerical typos

## Find the problem (when we don't know what it is)

- Make a new column and convert **amount** to numbers
- Find out where the conversion didn't work

```
filter(rivers, !is.na(amount), is.na(amount2))
```

```
## # A tibble: 1 x 5
##   river_name site      element amount amount2
##   <chr>      <chr>    <chr>   <chr>   <dbl>
## 1 Raquette  Mid stream Ca      <0.1    NA
```

- **is.na()** is **TRUE** when the value is missing (**NA**)
- **!** turns a **TRUE** into a **FALSE** (and vice versa)
- This asks, which values are not missing to begin with (**!is.na(amount)**) but *are* missing after the conversion (**is.na(amount2)**)

# Fixing numerical typos

## Find the problem (when we know what it is):

```
filter(rivers, amount == "<0.1")
```

```
## # A tibble: 1 x 5
```

```
##   river_name site      element amount amount2
```

```
##   <chr>      <chr>      <chr>  <chr>    <dbl>
```

```
## 1 Raquette  Mid stream Ca      <0.1      NA
```

# Fixing numerical typos

## Find the problem (when we know what it is):

```
filter(rivers, amount == "<0.1")
```

```
## # A tibble: 1 x 5
```

```
##   river_name site      element amount amount2
```

```
##   <chr>      <chr>      <chr>  <chr>    <dbl>
```

```
## 1 Raquette  Mid stream Ca      <0.1      NA
```

## Fix problem

```
rivers <- mutate(rivers, amount = if_else(amount == "<0.1", "0", amount))
```

# Fixing numerical typos

## Find the problem (when we know what it is):

```
filter(rivers, amount == "<0.1")
```

```
## # A tibble: 1 x 5
```

```
##   river_name site      element amount amount2
```

```
##   <chr>      <chr>      <chr>  <chr>    <dbl>
```

```
## 1 Raquette  Mid stream Ca      <0.1      NA
```

## Fix problem

```
rivers <- mutate(rivers, amount = if_else(amount == "<0.1", "0", amount))
```

## Correct the class

```
rivers <- mutate(rivers, amount = as.numeric(amount))
```

# Put it together...

```
rivers <- read_csv("./data/rivers_correct.csv") %>%  
  clean_names() %>%  
  rename(element = ele, amount = amo) %>%  
  select(-wea) %>%  
  mutate(river_name = case_when(river_name %in% c("Grase", "Grass", "grasse") ~ "Grasse",  
                                river_name == "raquette" ~ "Raquette",  
                                TRUE ~ river_name),  
         site = case_when(site == "Dow stream" ~ "Down stream",  
                           site == "Upstream" ~ "Up stream",  
                           TRUE ~ site),  
         amount = if_else(amount == "<0.1", "0", amount),  
         amount = as.numeric(amount))
```

## And you have a clean, corrected data frame ready to use

- You have not changed the original data
- You have a record of all corrections
- You can alter these corrections at any time
- You have formatted your data for use in R

# Dates and Times

(Or why does R hate me?)

# Dates and Times

- Date/times aren't always recognized as date/times

```
geolocators <- read_csv("./data/geolocators.csv", col_names = c("time", "light"))
geolocators
```

```
## # A tibble: 21 x 2
##   time                light
##   <chr>              <dbl>
## 1 02/05/11 22:29:59      64
## 2 02/05/11 22:31:59      64
## 3 02/05/11 22:33:59      38
## 4 02/05/11 22:35:59      38
## 5 02/05/11 22:37:59      34
## 6 02/05/11 22:39:59      30
## # ... with 15 more rows
```

Here **time** column is considered **chr** (character/text)





# lubridate package

- Part of **tidyverse**, but needs to be loaded
- Great for converting date/time formats

```
library(lubridate)
geolocators <- mutate(geolocators, time = dmy_hms(time))
geolocators
```

```
## # A tibble: 21 x 2
##   time                light
##   <dtm>              <dbl>
## 1 2011-05-02 22:29:59     64
## 2 2011-05-02 22:31:59     64
## 3 2011-05-02 22:33:59     38
## 4 2011-05-02 22:35:59     38
## 5 2011-05-02 22:37:59     34
## 6 2011-05-02 22:39:59     30
## # ... with 15 more rows
```

# lubridate package

date/time	function	class
2018-01-01 13:09:11	<code>ymd_hms()</code>	dtm (POSIXct/POSIXt)
12/20/2019 10:00 PM	<code>mdy_hm()</code>	dtm (POSIXct/POSIXt)
31/01/2000 10 AM	<code>dmy_h()</code>	dtm (POSIXct/POSIXt)
31-01/2000	<code>dmy()</code>	Date

**lubridate** is smart enough to detect AMs and PMs

## 5. Saving data

# Saving data

## Keep yourself organized

- Keep your R-created data in a **different** folder from your 'raw' data
- If you have a lot going on, split your work into several scripts, and number the data sets produced:
  - **1\_cleaned.csv**
  - **2\_summarized.csv**
  - **3\_graphing.csv**

# Saving data

## Keep yourself organized

- Keep your R-created data in a **different** folder from your 'raw' data
- If you have a lot going on, split your work into several scripts, and number the data sets produced:
  - **1\_cleaned.csv**
  - **2\_summarized.csv**
  - **3\_graphing.csv**

## Save your data to file:

```
write_csv(rivers, "../Datasets/rivers_cleaned.csv")
```

# Dealing with data

## 1. Loading data

- Get your data into R

## 2. Looking for problems

- Typos
- Incorrectly loaded data

## 3. Fixing problems

- Corrections
- Renaming

## 4. Setting formats

- Dates
- Numbers
- Factors

## 5. Saving your data

# Wrapping up: Common mistakes

## Forgetting to use `as.character()` when switching from factor to numeric

- Applies especially if you use `read.csv()` (R base function which often creates factors)
- To convert factor to numeric use: `as.numeric(as.character(my_factor))`



# Wrapping up: Common mistakes

## Forgetting to use `as.character()` when switching from factor to numeric

- Applies especially if you use `read.csv()` (R base function which often creates factors)
- To convert factor to numeric use: `as.numeric(as.character(my_factor))`

## Assuming your data is in one format when it's not

- Print your data to the console and use `skim()` to explore the format of your data
- Use `skim()`, `count()`, `filter()`, `select()`, `ggplot()` to explore the content of your data

# Wrapping up: Common mistakes

## Confusing pipes with function arguments

- Pipes (`%>%`) pass the *output* from one function as *input* to the next function:

```
my_data <- my_data %>%           # Pass my_data
  filter(my_column > 5) %>%      # Pass my_data, filtered
  select(my_column, my_second_column)
```

- Arguments may be on different lines, but all part of *one* function

```
my_data <- my_data %>%           # Pass my_data
  mutate(my_column1 = replace(...), # No passing (no pipes!)
         my_column2 = replace(...), # Instead, give 3 arguments to mutate:
         my_column3 = replace(...)) # Arguments separated by ",", and surrounded by ( )
```

# Wrapping up: Further reading

- R for Data Science
  - [Chapter 5: Transforming data](#)
  - [Chapter 8: RStudio Projects](#)
  - [Chapter 14: Strings](#)
  - [Chapter 15: Factors](#)
  - [Chapter 18: Pipes](#)